

MarsLab: A Martian Rover Simulator for Planetary Rover Autonomous Navigation

Hoyun Kim^{1*}, Beomsu Kim¹, and Giseop Kim^{1†}

Abstract—Future Mars missions will require rover autonomy that can operate across unstructured terrain, changing illumination, atmospheric dust, and limited communication. Simulation is a practical way to study these conditions before deployment, but existing Mars-relevant resources differ in scope, including mission-oriented simulators, fixed analog datasets, task-specific environments, and open robotics interfaces. In this context, we present MarsLab, an open-source, ROS2-native Mars rover simulator for autonomy and navigation algorithm development. MarsLab combines HiRISE-derived and procedural terrain with customizable rock, crater, solar-illumination, and atmospheric-dust settings, and runs a Perseverance-class rover model in NVIDIA Isaac Sim. The runtime publishes RGB, depth, RGB-D point clouds, LiDAR, IMU, wheel odometry, and Ground Truth (GT) pose data through standard ROS2 topics. We demonstrate MarsLab with Simultaneous Localization and Mapping (SLAM) benchmarks across sensing modalities, dust levels, scene geometry, and route length, and with Visual Place Recognition (VPR) benchmarks over repeated Mars Base traversals under illumination and dust changes. The results illustrate how controlled scene variation and shared GT trajectories can be used to compare trajectory-level estimation and image-level place recognition within the same simulator. Our Project Page: <https://kimhoyun-robotair.github.io/MarsLab/>.

I. INTRODUCTION

Mars exploration is entering a renewed robotic phase. In 2026, NASA approved implementation of its support to ESA’s Rosalind Franklin rover and selected SpaceX Falcon Heavy for the mission’s late-2028 launch opportunity [1]. NASA also announced Space Reactor-1 Freedom, a nuclear-electric spacecraft intended to reach Mars before the end of 2028 and deploy a Skyfall payload of Ingenuity-class helicopters [2]. Together with Perseverance AutoNav and Ingenuity, these plans point toward Mars missions that combine rover mobility, aerial scouting, commercial launch, and increasingly autonomous robotic decision making [3], [4].

Mars exploration remains difficult despite this momentum. Mission systems have progressed from Mars Exploration Rover visual odometry [5] to Perseverance driving autonomy [3], orbital-map localization [6], and coupled SLAM/navigation pipelines [7]. Yet surface autonomy must operate without Global Navigation Satellite System (GNSS), with delayed communication, and over low-texture regolith,

¹H. Kim, B. Kim, and G. Kim are with the Department of Robotics and Mechatronics Engineering, DGIST, Daegu, Republic of Korea {hoyunkim, beomsu.kim, gsk}@dgist.ac.kr.

*First author: Hoyun Kim. †Corresponding author: Giseop Kim.

This work was supported by the National Research Foundation of Korea (NRF) grants funded by the Korea government (MSIT) (No. RS-2026-25492530 and No. RS-2026-25517444), and by the InnoCORE program of the Ministry of Science and ICT (No. 26-InnoCORE-01).




Fig. 1: **MarsLab overview.** MarsLab supports Mars-rover autonomy and navigation algorithm development through HiRISE-based customizable scenes and environment variables, together with sensor topic streams. Users can vary terrain, rocks, craters, solar illumination, and atmospheric dust (Fig. 4–Fig. 6), subscribe to RGB, depth/RGB-D, and 3D LiDAR sensor topics from the rover runtime, and evaluate autonomy modules including SLAM and visual place recognition, as demonstrated in our experiments.

repetitive terrain, changing illumination, dust, and tight resource budgets [8]. Localization is especially limiting: pose error propagates into mapping and planning, while repeatable ground-truth pose data on Mars remains scarce.

Because these developments cannot be reproduced through field testing on Mars, simulation-based validation is essential for evaluating autonomy stacks before deployment. However, as summarized in Table I, existing resources rarely combine openness with full-stack autonomy and navigation support: mission simulators are typically not released as community tools [9], [10], several prior environments emphasize narrower task settings such as SLAM, photorealism,

TABLE I

Concise survey of Mars rover simulation platforms and benchmark support.

Column abbreviations and symbols: **gen.** = generation; **seqs.** = sequences; **sens.** = sensor modalities; ✓ = reported/core capability; ○ = partial, indirect, or not primary; ✗ = not reported; **D** = depth; **L** = 3D LiDAR; **seg.** = semantic or instance labels; **NR** = not reported.

Simulator	Year	Scene gen.	Engine	Goal	Scenes/Seqs.	Rock	Dust	Light	Sens.	ROS2	Open
ROAMS [9]	2004	Mars-like terrain & mission simulation	Custom / JPL DARTS-DSHELL	Rover mobility, dynamics, onboard SW V&V	Mission/testbed scenarios; Monte Carlo possible	✓	✗	○	Stereo cameras; IMU; encoders; sun sensors	✗	○
ENav [10]	2020	Mars 2020 ENav test terrains	ROS-based + HDSim/RSVP	Autonomous nav V&V	Mars 2020 test scenarios / Monte Carlo runs	✓	✗	○	Stereo/NavCam point cloud; camera rendering	✗	✗
Giubilato et al. [11]	2020	Simulated Martian / planetary-like environment	ROS/Gazebo	Visual & LiDAR SLAM evaluation	Long/Short simulated rover sequences / datasets	✓	✗	○	Cam / 3D L; stereo + LiDAR seqs.	✗	✓
MarsSim [12]	2023	Multiscale Mars terrain simulation	ROS/Gazebo	High-fidelity physical & visual rover simulation	Pahrump Hills-like and generated scenes; released seqs. NR	✓	○	✓	RGB/visual; locomotion data; D/L NR	✗	✗
ISMRS [13]	2024	Digital-twin and asset-based scenes	Isaac Sim	Mars rover simulation + ML data synthesis	Real-world recordings + synthetic data	✓	✗	✓	RGB/stereo; 2D LiDAR; IMU; wheel odom; seg.	✓	✗
RLRoverLab [14]	2024	Synthetic rover training scenes	Isaac Lab	RL navigation / manipulation / control	RL task scenes	○	✗	○	Task-dependent; RGB cam, height scan, Isaac obs.	✗	✓
MarsLab (Ours)	2026	HiRISE-based + customizable	Isaac Sim ROS2	Rover autonomy algorithm development	Customizable multi-scenes; SLAM/VPR/nav tests	✓	✓	✓	RGB/D/L; wheel odom; GT logs	✓	✓

Reinforcement Learning (RL), or synthetic data [11]–[14], and analog datasets provide fixed recordings [15].

To address this gap, we present **MarsLab**: an open-source, ROS2 [16]-native Mars rover simulator for repeatable autonomy evaluation across customizable terrain, lighting, dust, sensing, rover motion, and time-aligned ground truth pose data. MarsLab supports SLAM and navigation algorithm development through HiRISE-based customizable scenes [17] and standard rover sensor topics (Fig. 1).

Our contributions are:

- 1) **MarsLab**, an open-source, ROS2-native Mars rover simulation and validation environment built on NVIDIA Isaac Sim;
- 2) a scientifically grounded scene-generation pipeline that couples HiRISE-derived terrain models with procedural Mars-like rocks, craters, illumination, and atmospheric dust to produce Mars-relevant environmental variation;
- 3) a reproducible, multi-scene benchmark for SLAM, localization, place recognition, planning, and navigation, demonstrated through SLAM and VPR experiments.

II. RELATED WORK

A. Planetary Simulation and Autonomy Validation

Simulation is an established approach for planetary-surface validation, but existing Mars tools remain limited in scope and availability. ROAMS and the Mars 2020 ENav simulator are mission-oriented and not community benchmarks [9], [10]. ROS/Gazebo and Isaac Sim systems such as Giubilato et al. [11], MarsSim [12], ISMRS [13], and RLRoverLab [14] address SLAM evaluation, visual/physical fidelity, digital twins, or RL, but do not jointly provide an open ROS2 benchmark with customizable Mars terrain,

illumination, dust, sensors, and ground truth poses. Table I positions MarsLab against this landscape.

B. Planetary Datasets, Semantics, and Traversability

Planetary datasets capture important perceptual information, but their recorded trajectories, environmental conditions, and sensor configurations are typically fixed after release. AI4MARS supplies terrain labels for Mars imagery [19]. The Canadian energy-aware navigation dataset provides rover-collected multi-sensor logs [15], and DLR S3LI provides handheld stereo, solid-state LiDAR, and inertial recordings from a planetary-analog environment [20]. They support offline perception studies, but cannot vary GT pose trajectories, terrain statistics, dust, illumination, or sensor suites to measure closed-loop effects.

C. Long-Term Localization and Mars SLAM

Long-term localization is central to planetary autonomy because drift propagates to mapping, traversability assessment, and planning. The Mars Exploration Rovers have used visual odometry [5]. Orbital-map matching for global correction has been evaluated in terrestrial analog terrain [7]. Censile has demonstrated onboard global localization on Perseverance in shadow mode [6]. Direct validation on Mars remains limited by cost, sparse ground truth pose data, and non-repeatable conditions.

MarsLab addresses this gap with HiRISE-derived and procedural terrain, a Perseverance-class rover, controlled illumination and dust, synchronized ground truth poses, and a standardized ROS2 benchmark.

III. MARSLAB ARCHITECTURE

A. System Overview

MarsLab is an open-source, ROS2-native simulator built on NVIDIA Isaac Sim 5.1 [21]. As shown in Fig. 2, the




Fig. 2: **MarsLab pipeline.** MarsLab converts planetary terrain, surface assets, and rover models into customizable Mars experiments. The pipeline builds scene variants with controlled terrain generation, rocks, craters, illumination, and dust; executes the rover in an Isaac Sim/ROS2 runtime under Mars gravity; and exports sensor, command, and GT pose for SLAM, localization, place recognition, planning, and navigation evaluation.

pipeline starts from reusable terrain, surface, and rover assets; constructs customizable Mars scenes; runs them with Isaac Sim physics and ROS2 streams; and exports experiment data for evaluation. Each run is specified by a Universal Scene Description (USD) scene, YAML runtime configuration, seed, and trajectory/ground-truth pose data. Runtime parameters are validated before launch, and randomized stages such as crater sampling, rock placement, and sensor noise share the explicit seed for reproducible replay.

The architecture is organized as an end-to-end autonomy-evaluation pipeline. The scene-construction layer builds Mars-relevant worlds from terrain sources and customizable surface and environmental parameters. The rover-runtime layer loads each world into Isaac Sim, applies Mars gravity, and provides rover dynamics, command topics, sensor streams, and coordinate frames. The benchmark layer exports GT trajectory and sensor data for evaluating SLAM, localization, navigation, VPR, and related autonomy algorithms. This separation lets users vary scene conditions without changing the rover interface or the evaluation format.

B. Terrain and Scene Generation

MarsLab combines real and procedural terrain sources. HiRISE Digital Elevation Model (DEM)s and co-registered orthomosaics provide real terrain geometry and texture information, with Mastcam-Z imagery as a surface-appearance reference [17], [22]. Procedural terrains are synthesized with `synthterrain` [23] and augmented with the rock and crater models below. The reference scenes span plains, craters, canyons, and habitat sites¹ (Fig. 3; Table II), while Fig. 4–6 show the customizable terrain, illumination, and dust factors used as controlled evaluation conditions.

¹“Mars Base” by MOJackal, via Sketchfab, <https://sketchfab.com/3d-models/mars-base-2bc6660c63b24158bddd20ff5d2235e>, licensed under CC BY 4.0.

Scene construction follows a deterministic asset-build process. Given a crop region, terrain scale, and generation seed, the builder places rocks, crater fields, rover assets, lighting, and atmospheric settings into a single world description. The generated scene is stored separately from the runtime configuration, so it can be replayed with different sensors or controllers and the same GT trajectory can be evaluated under modified lighting or dust. This design keeps MarsLab from being a fixed dataset: it is a controlled generator of related experiments whose differences are known before evaluation.

C. Rock Distribution

Rock populations follow the Golombek–Huertas cumulative fractional-area model [24]:

$$F_k(D) = k e^{-q(k)D}, \quad q(k) = 1.79 + \frac{0.152}{k}, \quad (1)$$

where D is rock diameter, k is the total rock-covered area fraction, and $F_k(D)$ is the cumulative area fraction covered by rocks with diameter at least D . Sampled diameters are used to scale USD rock assets² on the DEM surface, so the generated rock field follows the target size–frequency statistics rather than arbitrary placement (Fig. 4a).

D. Crater Distribution

Crater diameters are sampled from the Hartmann–Neukum Mars production function [25]. For the valid $D \geq 10$ m regime, MarsLab converts the cumulative SFD $C(D)$ into a normalized sampling CDF:

$$F(D) = 1 - \frac{C(D)}{C(D_{\min})}, \quad D \in [D_{\min}, D_{\max}]. \quad (2)$$

²“Mars Rocks” by Ivan Vakulko (milos4), via Sketchfab, <https://sketchfab.com/3d-models/mars-rocks-9f5c946255a24f1cb630ea95dceea587>, licensed under CC BY 4.0.




Fig. 3: **Reference MarsLab scenes.** MarsLab can generate additional terrain instances, and these four reference scene pairs are provided as scenes for evaluation. The green rectangle in each upper image indicates the approximate area represented by the corresponding lower scene: (a) a plain scene from Jezero crater, (b) a main-crater scene from Arcadia Planitia, (c) a Mars Base scene that augments the Jezero terrain with a habitat asset, and (d) a Grand Canyon scene from Elysium Planitia. All scenes are built on HiRISE-derived terrain, providing diverse geometry, texture, and structure for SLAM and place recognition stress testing.

TABLE II

MarsLab reference scenes. Compact summary of the four evaluation scenes in Fig. 3. Rock distribution, illumination, and atmospheric dust denote user-customizable scene parameters, including rock size–frequency/spatial-density settings, solar azimuth/elevation, and dust optical depth.

Fig.	Reference scene	Scene size	Rock	Dust	Illumination	Target Use
(a)	Jezero plain/rim	0.5×0.5 km	✓	✓	✓	Navigation, terrain following, and low-texture localization baselines.
(b)	Main Crater	3.6×3.6 km	✓	✓	✓	Repetitive-terrain failure analysis and rocky slope traversal.
(c)	Mars Base	1.0×1.0 km	✓	✓	✓	Landmark- and feature-rich localization and VPR benchmarking.
(d)	Grand Canyon	3.0×3.0 km	✓	✓	✓	Long-range drift evaluation over extended canyon terrain with strong elevation changes and complex geometry.

Optional resurfacing masks follow crater-population clustering analysis [26], and diameter-dependent profiles use published fresh-crater morphometry [27]; thus crater density and shape remain tied to planetary chronology models (Fig. 4b).

E. Solar Illumination and Dust Environment

Illumination is treated as a customizable experimental variable. Direct surface irradiance follows Beer’s law with Mars dust optical depth τ :

$$I(z) = \begin{cases} I_0 e^{-\tau m(z)}, & z < 90^\circ, \\ 0, & z \geq 90^\circ, \end{cases} \quad (3)$$

where z is solar zenith angle, $I_0 \approx 589 \text{ W m}^{-2}$, and $m(z)$ is the Appelbaum–Flood relative optical air mass [28]. Diffuse sky illumination follows a τ -indexed COMIMART reduction [29]. The same trajectory can then be replayed with customizable solar azimuth/elevation and dust optical depth: sun changes redistribute shadows (Fig. 5), while increasing τ reduces contrast and visibility (Fig. 6) without changing geometry. By varying one environmental parameter at a time, such as sun position or dust optical depth, MarsLab enables vision-based algorithms to be tested under diverse conditions within the same scene.





Fig. 4: **Procedurally generated rock and crater scenes.** Rendered MarsLab terrain showing (a) the rock size–frequency distribution (Eq. 1) and (b) crater density and morphometry (Eq. 2).



(a) Azimuth 70° , elevation 13° (b) Azimuth 360° , elevation 13°

Fig. 5: **Controllable solar illumination.** The sun azimuth and elevation are explicit experimental variables; changing the solar position redistributes cast shadows and shading across the same terrain, directly affecting feature visibility for vision-based navigation algorithms.

TABLE III

Customizable MarsLab scene generation. Beyond the reference scenes, MarsLab lets users build HiRISE-based scenes by controlling terrain extent, surface clutter, crater fields, illumination, and atmospheric dust. The table summarizes the scene-generation factors exposed as user-controlled variables.

Stage	User-provided input	Customizable variation
Terrain base	HiRISE DEM/orthomosaic	Crop region, terrain scale, and generation seed.
Surface clutter	Rock assets and rock size–frequency parameters	Rock density, diameter range, spatial sampling, and placement seed.
Crater field	Mars crater production-function parameters	Crater density, diameter range, morphology, and resurfacing masks.
Illumination	Solar-position and irradiance configuration	Sun azimuth, elevation, shadow direction, and scene shading.
Atmosphere	Mars dust optical-depth setting	Dust optical depth τ , sky brightness, and image contrast.




Fig. 6: **Atmospheric-dust sweep.** MarsLab uses the optical-depth parameter τ in Eq. 3 as a controlled evaluation condition for vision-based algorithms: increasing τ brightens the sky and ground, lowers image contrast, and reduces visibility while leaving the terrain geometry unchanged.

F. Rover Runtime and Physics

The runtime uses a Perseverance (M2020) model from NASA-JPL’s public URDF release [18], converted to USD and simulated under Mars gravity (3.72 m/s^2). A sensor graph and `roclpy` bridge subscribe to `cmd_vel` and publish odometry, IMU, RGB, depth, RGB-D point cloud, 2D/3D LiDAR, TF, robot description, and ground-truth poses in REP-103/REP-105 frames (Fig. 7). GT trajectories are written in TUM format and replayed by a pure-pursuit follower with bounded commands ($v \leq 0.5 \text{ m/s}$, $|\omega| \leq 0.2 \text{ rad/s}$, 50 Hz), giving a reference data stream for Absolute Trajectory Error (ATE)/RPE evaluation. The resulting sensor streams and GT trajectory data provide the common interface used by the runtime and the benchmarks below. The runtime also supports experiments and analyses under varying sensor-noise

```

marslab@ros2:~$ ros2 topic list
/rover/cmd_vel[subscribe]      /rover/imu[publish]
/clock[publish]                /rover/imu_noisy[publish]
/tf[publish]                   /rover/rgb/image_raw[publish]
/tf_static[publish]            /rover/rgb/camera_info[publish]
/rover/joint_states[publish]    /rover/depth/image_raw[publish]
/rover/robot_description[publish] /rover/depth/points[publish]
/rover/odom[publish]           /rover/lidar/points[publish]
/rover/GT_Trajectory[publish]  /rover/scan[publish]

```

Fig. 7: **MarsLab ROS2 topic interface.** Default subscribed and published topics exposed by the rover runtime.

TABLE IV

ATE RMSE (m) across MarsLab scenes and SLAM stacks. Camera methods are run at low ($\tau=0.5$) and high ($\tau=6.0$) dust; LiDAR is run once per scene. Per-row best in **bold**; **X** marks a tracking failure.

Scene	ORB-SLAM		RTAB-Map		MOLA
	$\tau=0.5$	$\tau=6.0$	$\tau=0.5$	$\tau=6.0$	LiDAR
Mars Base	0.59	4.70	0.48	0.64	0.20
Main Crater Interior	40.28	X	8.84	10.61	0.13
Grand Canyon	X	X	12.28	19.08	0.38

X = monocular tracking failure (no usable trajectory).

conditions by providing zero-mean Gaussian perturbations with customizable standard deviations for IMU and wheel-odometry measurements, as well as optional depth-camera post-processing with customizable noise mean and scale.

IV. EXPERIMENTS

We evaluate MarsLab with two studies: a multi-scene SLAM benchmark comparing RGB, RGB-D/wheel-odometry, and 3D-LiDAR methods under controlled dust, and a VPR benchmark over repeated *Mars Base* traversals.

A. SLAM Benchmark

Setup. The SLAM benchmark uses three representative scenes: *Mars Base* for flat terrain with structural landmarks, *Main Crater Interior* for rock-scattered crater terrain, and *Grand Canyon* for large elevation changes and occlusion. We evaluate ORB-SLAM [30] on monocular RGB, RTAB-Map [31] on RGB-D plus wheel odometry, and MOLA [32] on 3D LiDAR. ORB-SLAM and RTAB-Map are run at $\tau=0.5$ and $\tau=6.0$; MOLA is run once per scene because the dust model affects radiance but not range.

Protocol. For each scene, we design the GT trajectory to keep the rover on traversable terrain by accounting for rock distribution and local terrain slope. The pure-pursuit follower executes the trajectory, and MarsLab logs ground truth poses in TUM format. We report ATE RMSE after Umeyama SE(3) alignment. Table V and Fig. 10 define the GT trajectories, Fig. 8 overlays trajectories, and Fig. 9 visualizes reconstructed maps. All runs use Isaac Sim 5.1, ROS2 Jazzy, and an NVIDIA RTX 5070 Ti workstation.

Results. Table IV shows a clear modality ordering. MOLA remains sub-metre on all scenes (0.13–0.38 m), RTAB-Map ranges from 0.48 m on *Mars Base* to 8–19 m on rugged scenes, and monocular ORB-SLAM is fragile, reaching 40.3 m on *Main Crater Interior* and losing tracking on *Grand Canyon*. Increasing dust from $\tau=0.5$ to $\tau=6.0$ mainly affects vision: ORB-SLAM grows $7.9\times$ on *Mars Base* and fails on *Main Crater Interior*, while RTAB-Map degrades more gradually because depth and wheel odometry anchor scale and motion. MOLA is unchanged, consistent with a dust model that attenuates radiance but not range.

Failure modes. Scene geometry also matters. *Mars Base* is easiest for vision, whereas canyon walls and self-similar crater terrain starve monocular features, causing the tracking losses marked in Fig. 8 and Table IV. MOLA’s error instead follows GT trajectory length, rising from 0.13 m on the 393 m




Fig. 8: **Multi-scene SLAM benchmark: estimated vs. ground-truth trajectories.** Top-down view of every run—the estimated trajectory (red) versus ground truth (black) after Umeyama alignment, over the reconstructed point-cloud map (coloured by height); rows are scenes, columns are sensing modality and dust optical depth. Grey tiles (X) mark monocular ORB-SLAM runs that lost tracking and produced no usable trajectory. Scene names and GT trajectory lengths are listed in Table V; GT trajectory geometry is shown in Fig. 10; per-run ATE in Table IV; reconstructed 3D maps in Fig. 9.




Fig. 9: **Reconstructed 3D maps: Visual SLAM vs. LiDAR SLAM.** Point-cloud maps reconstructed for each scene, comparing visual SLAM (top) against LiDAR SLAM (bottom; MOLA). For (a) *Mars Base* the visual-SLAM map is from ORB-SLAM (RGB); for (b) *Main Crater Interior* and (c) *Grand Canyon* it is from RTAB-Map (RGB-D+Odom), whose depth and wheel-odometry inputs remained stable where monocular ORB-SLAM lost tracking. Points are coloured by elevation.




Fig. 10: **GT trajectories.** Compact overview of the three indexed ground-truth trajectories used in the SLAM benchmark.

TABLE V

SLAM benchmark GT trajectory metadata. Each GT trajectory is authored in a MarsLab scene, replayed by the rover runtime, and logged from Isaac Sim pose streams in TUM format for ATE/RPE evaluation.

Scene	Length	GT Trajectory Role
Mars Base	≈ 440 m	Feature- and landmark-rich trajectory for testing localization in structured habitat scenes.
Main Crater Interior	≈ 393 m	Rock-rich trajectory for testing localization under repetitive terrain and sparse visual features.
Grand Canyon	≈ 648 m	Long-distance trajectory for testing drift across extended canyon terrain with large elevation changes.

Main Crater Interior trajectory to 0.38 m on the 648 m *Grand Canyon* trajectory, or only 0.03–0.06% of distance travelled. This separation of dust, geometry, and GT trajectory length illustrates the controlled attribution MarsLab is designed to support.

B. Visual Place Recognition

Setup. Beyond trajectory-level SLAM, we evaluate VPR as image retrieval over repeated *Mars Base* traversals. We compare NetVLAD [33], AnyLoc [34], and BoQ [35] on 4456 query and 4455 database frames recorded every 0.1 m. A retrieval is correct when one of the top- N database images lies within 2 m of the query pose; we report Recall@ N for $N \in \{1, 5, 10\}$.

Protocol. We test two appearance changes on the same Mars Base route. *Day/Dark* matches day-lit queries to a dark database at $\tau=0.05$, while *Dust* matches low-dust daytime queries at $\tau=0.05$ to high-dust daytime database images at $\tau=0.60$.

Results. Table VI shows that BoQ performs best in both protocols, reaching 91.44% R@1 for Day/Dark and 94.65% R@1 for Dust. NetVLAD ranks second across all metrics, and AnyLoc recovers by R@10 despite lower R@1 accuracy. Fig. 11 confirms the same trend qualitatively: the top-1 matches for the habitat dome and lander/Starship still depict the correct places under viewpoint and appearance changes.

C. Experiment Summary

Taken together, the experiments show how MarsLab supports autonomy benchmarks across both trajectory-level estimation and image-level place recognition. The SLAM benchmark varies dust level, long-range GT trajectory, rock distribution, and scene structure, spanning feature- and landmark-




Fig. 11: **Visual place recognition: query vs. retrieved.** Qualitative examples for the two Table VI protocols: (i) day-lit queries matched to a dark database at $\tau=0.05$ and (ii) low-dust daytime queries at $\tau=0.05$ matched to high-dust daytime database images at $\tau=0.60$. Each pair shows the query frame and the top-1 retrieved database frame (R@1).

TABLE VI

Visual place recognition recall on the MarsLab Mars-base benchmark using the strict 2 m positive radius. Results are Recall@ N (%). Best values in each protocol/metric are shown in bold.

Method	Day/Dark, $\tau=0.05$			Dust, $\tau=0.05 \rightarrow 0.60$		
	R@1	R@5	R@10	R@1	R@5	R@10
NetVLAD [33]	77.90	95.35	98.36	85.02	98.18	99.33
AnyLoc [34]	64.04	92.21	97.17	74.89	95.71	99.06
BoQ [35]	91.44	99.19	99.80	94.65	99.37	99.78

poor terrain as well as landmark-rich *Mars Base* conditions. It also exposes the sensing modalities needed for practical rover navigation: monocular camera, RGB-D, LiDAR, and wheel odometry, with IMU streams available for additional tests even though this paper does not include an IMU-based SLAM run. The VPR benchmark complements this by using the same *Mars Base* route to change dust and illumination while preserving pose labels, allowing users to design controlled VPR protocols and compare retrieval behavior across methods. The different R@1 scores in Table VI confirm that these appearance changes produce meaningful algorithm-dependent differences rather than trivial retrieval cases.

V. DISCUSSION

MarsLab combines customizable Mars environments with ROS2 sensor interfaces and ground-truth outputs, enabling repeatable autonomy experiments under different terrain, illumination, and dust settings. Future work will extend the IMU noise model to include constant biases and time-varying drift, and incorporate particle-based terramechanics for studying deformable regolith and wheel–soil interaction.

In addition to these modeling extensions, future work will introduce multi-robot and heterogeneous platform support, including different rover types and aerial scouts with platform-specific dynamics and controllers. These extensions will broaden evaluation beyond SLAM and VPR to computer-vision tasks, traversability estimation, path planning, and multi-robot cooperation, using shared Mars environments, customizable sensor interfaces, and per-robot ground-truth outputs.

VI. CONCLUSION

We presented MarsLab, an open-source, ROS2-native Mars rover simulator for reproducible autonomy evaluation. MarsLab combines HiRISE-derived and procedural terrain with distributed rock, crater, illumination, and dust models, then plays each scene and GT trajectory through an Isaac Sim/ROS2 runtime with ground truth pose data. The experiments show that MarsLab lets users vary sensing modality, dust, scene geometry, and GT trajectory length while keeping the remaining evaluation conditions explicit. By coupling Mars-relevant scene generation with rover motion and standard robotics interfaces, MarsLab provides a practical testbed for localization, mapping, navigation, place recognition, and broader rover-autonomy algorithm development.

REFERENCES

- [1] NASA Science Editorial Team, “NASA Begins Implementation for ESA’s Rosalind Franklin Mission to Mars,” Apr. 2026, NASA Science, Apr. 16, 2026.
- [2] L. E. Low, “NASA Unveils Initiatives to Achieve America’s National Space Policy,” Mar. 2026, NASA, Release 26-025, Mar. 24, 2026.
- [3] V. Verma *et al.*, “Autonomous robotics is driving Perseverance rover’s progress on Mars,” *Science Robotics*, vol. 8, no. 80, p. eadi3099, 2023.
- [4] NASA Jet Propulsion Laboratory, “NASA Uses Two Worlds to Test Future Mars Helicopter Designs,” Nov. 2023.
- [5] M. Maimone *et al.*, “Two years of Visual Odometry on the Mars Exploration Rovers,” *Journal of Field Robotics*, vol. 24, no. 3, pp. 169–186, 2007.
- [6] J. Nash *et al.*, “Censible: A robust and practical global localization framework for planetary surface missions,” in *2024 IEEE International Conference on Robotics and Automation (ICRA)*, 2024, pp. 8642–8648.
- [7] D. Geromichalos *et al.*, “SLAM for autonomous planetary rovers with global localization,” *Journal of Field Robotics*, vol. 37, no. 5, pp. 830–847, 2020.
- [8] Y. Gao and S. Chien, “Review on space robotics: Toward top-level science through space exploration,” *Science Robotics*, vol. 2, no. 7, p. eaan5074, 2017.
- [9] A. Jain *et al.*, “Recent developments in the ROAMS planetary rover simulation environment,” in *2004 IEEE Aerospace Conference Proceedings (IEEE Cat. No. 04TH8720)*, vol. 2, 2004, pp. 861–876 Vol.2.
- [10] O. Toupet *et al.*, “A ROS-based simulator for testing the enhanced autonomous navigation of the Mars 2020 rover,” in *2020 IEEE Aerospace Conference*, 2020, pp. 1–11.
- [11] R. Giubilato *et al.*, “Simulation framework for mobile robots in planetary-like environments,” in *2020 IEEE 7th International Workshop on Metrology for AeroSpace (MetroAeroSpace)*, 2020, pp. 594–599.
- [12] R. Zhou *et al.*, “MarsSim: A high-fidelity physical and visual simulation for Mars rovers,” *IEEE Transactions on Aerospace and Electronic Systems*, vol. 59, no. 2, pp. 1879–1892, 2023.
- [13] G. Wan *et al.*, “High fidelity Mars rover simulation platform based on digital twin and machine learning integration: Innovation and application of ISMRS,” 08 2024.
- [14] A. B. Mortensen and S. Bøgh, “RLRoverLAB: An advanced reinforcement learning suite for planetary rover simulation and training,” in *2024 International Conference on Space Robotics (iSpaRo)*, 2024, pp. 273–277.
- [15] O. Lamarre *et al.*, “The Canadian Planetary Emulation Terrain Energy-Aware Rover Navigation Dataset,” *The International Journal of Robotics Research*, 2020.
- [16] Open Robotics, “ROS: Robot operating system,” <https://www.ros.org/>, accessed: Jun. 24, 2026.
- [17] A. S. McEwen *et al.*, “Mars Reconnaissance Orbiter’s High Resolution Imaging Science Experiment (HiRISE),” *Journal of Geophysical Research: Planets*, vol. 112, no. E5, 2007.
- [18] NASA Jet Propulsion Laboratory, “m2020-urdf-models: Perseverance and Ingenuity URDF models,” 2022, public release URS307049, URS309682; credit NASA/JPL-Caltech; accessed 2026-05-29.
- [19] R. M. Swan *et al.*, “AI4MARS: A dataset for terrain-aware autonomous driving on Mars,” in *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2021, pp. 1982–1991.
- [20] R. Giubilato *et al.*, “Challenges of SLAM in extremely unstructured environments: The DLR planetary stereo, solid-state LiDAR, inertial dataset,” *IEEE Robotics and Automation Letters*, vol. 7, no. 4, pp. 8721–8728, 2022.
- [21] NVIDIA Corporation, “Isaac Sim: Robotics simulation and synthetic data generation,” <https://docs.isaacsim.omniverse.nvidia.com/latest/index.html>, accessed: Jun. 24, 2026.
- [22] A. G. Hayes *et al.*, “Pre-Flight Calibration of the Mars 2020 Rover Mastcam Zoom (Mastcam-Z) Multispectral, Stereoscopic Imager,” *Space Science Reviews*, vol. 217, no. 2, p. 29, 2021.
- [23] NASA Ames Research Center / NeoGeographyToolkit, “synthterrain: Tools for the generation of synthetic planetary terrain,” 2023, apache-2.0; accessed 2026-05-29.
- [24] M. P. Golombek *et al.*, “Size-frequency distributions of rocks on the northern plains of Mars with special reference to Phoenix landing surfaces,” *Journal of Geophysical Research: Planets*, vol. 113, no. E3, 2008.
- [25] W. K. Hartmann and G. Neukum, “Cratering chronology and the evolution of Mars,” *Space Science Reviews*, vol. 96, no. 1, pp. 165–194, 2001.
- [26] G. Michael *et al.*, “Planetary surface dating from crater size–frequency distribution measurements: Spatial randomness and clustering,” *Icarus*, vol. 218, no. 1, pp. 169–177, 2012.
- [27] W. A. Watters *et al.*, “Morphometry of small recent impact craters on Mars: Size and terrain dependence, short-term modification,” *Journal of Geophysical Research: Planets*, vol. 120, no. 2, pp. 226–254, 2015.
- [28] J. Appelbaum and D. J. Flood, “Solar radiation on Mars,” *Solar Energy*, vol. 45, no. 6, pp. 353–363, 1990.
- [29] Vicente-Retortillo, Álvaro *et al.*, “A model to calculate solar radiation fluxes on the Martian surface,” *J. Space Weather Space Clim.*, vol. 5, p. A33, 2015.
- [30] C. Campos *et al.*, “ORB-SLAM3: An accurate open-source library for visual, visual-inertial, and multimap SLAM,” *IEEE Transactions on Robotics*, vol. 37, no. 6, pp. 1874–1890, 2021.
- [31] M. Labbé and F. Michaud, “RTAB-Map as an open-source LiDAR and visual simultaneous localization and mapping library for large-scale and long-term online operation,” *Journal of Field Robotics*, vol. 36, no. 2, pp. 416–446, 2019.
- [32] J. L. Blanco-Claraco, “A Modular Optimization Framework for Localization and Mapping,” in *Proceedings of Robotics: Science and Systems*, Freiburg, Germany, June 2019.
- [33] R. Arandjelovic *et al.*, “NetVLAD: CNN architecture for weakly supervised place recognition,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2016.
- [34] N. Keetha *et al.*, “AnyLoc: Towards universal visual place recognition,” *IEEE Robotics and Automation Letters*, vol. 9, no. 2, pp. 1286–1293, 2024.
- [35] A. Ali-bey *et al.*, “BoQ: A place is worth a bag of learnable queries,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2024, pp. 17 794–17 803.